

# Notes on Generative Models: Adversarial Autoencoders and GANs

October 9, 2025

## 1 Adversarial Autoencoders (AAEs)

### 1.1 Overview

Adversarial Autoencoders (AAEs), introduced by Makhzani et al. (ICLR 2016), combine autoencoders with adversarial training to learn meaningful latent representations and generate data samples. They enforce a prior distribution on the latent space using a discriminator, inspired by Generative Adversarial Networks (GANs).

### 1.2 Core Components

- Encoder: Maps input  $x$  to latent code  $z$ , i.e.,  $q(z|x)$ .
- Decoder: Reconstructs input from  $z$ , i.e.,  $p(x|z)$ .
- Discriminator: Ensures the latent distribution  $q(z)$  matches a chosen prior  $p(z)$  (e.g., Gaussian).

### 1.3 Training Objective

The AAE optimizes two losses:

- Reconstruction Loss: Minimizes  $\|x - \hat{x}\|^2$  to ensure accurate input reconstruction.
- Adversarial Loss: Matches  $q(z)$  to  $p(z)$  using a GAN-like objective:

$$\min_q \max_D \mathbb{E}_{p(z)}[\log D(z)] + \mathbb{E}_{q(z)}[\log(1 - D(z))]$$

### 1.4 Advantages

- Structured latent space suitable for generative tasks.
- Flexibility to use any prior distribution (e.g., Gaussian, mixture models).
- Applications in data generation, semi-supervised learning, and representation learning.

## 2 Adversarial Autoencoders Paper (Makhzani et al., ICLR 2016)

### 2.1 Key Contributions

- Introduced AAEs as a probabilistic framework combining autoencoders and adversarial training.
- Replaced the KL-divergence penalty of VAEs with adversarial training for flexible prior matching.
- Demonstrated applications in generative modeling, semi-supervised learning, and representation learning.

### 2.2 Methodology

- Architecture: Encoder ( $q(z|x)$ ), decoder ( $p(x|z)$ ), and discriminator to match  $q(z)$  to  $p(z)$ .
- Training: Alternates between minimizing reconstruction loss and adversarial loss.
- Variants: Supervised AAEs (incorporate labels), semi-supervised AAEs, and denoising AAEs.

### 2.3 Experiments

- Datasets: MNIST, CIFAR-10, SVHN.
- Results: Generated sharp samples, achieved competitive semi-supervised classification, and learned meaningful latent representations.
- Comparisons: Outperformed VAEs in sample quality and flexibility, retained reconstruction unlike GANs.

### 2.4 Limitations

- Training instability due to adversarial component.
- Potential mode collapse in latent space.
- Higher computational cost than standard autoencoders.

## 3 VAEs vs. AAEs: Flexibility in Prior Matching

### 3.1 Statement

VAEs use KL-divergence to match the latent distribution  $q(z|x)$  to a simple prior (e.g.,  $\mathcal{N}(0, I)$ ), which is restrictive for complex distributions. AAEs use adversarial training, allowing any prior distribution.

### 3.2 Explanation

- VAEs: Optimize:

$$\mathcal{L}_{\text{VAE}} = \mathcal{L}_{\text{recon}}(x, \hat{x}) + \text{KL}(q(z|x)||p(z))$$

KL-divergence is tractable only for simple priors (e.g., Gaussian), limiting latent space complexity.

- AAEs: Use a discriminator to match  $q(z)$  to any  $p(z)$ , enabling complex priors like mixture models without computing KL-divergence.

### 3.3 Example: MNIST Digits

- VAE: Forces all digit latent codes into a single Gaussian, leading to blurry samples and poor separation of digit classes.
- AAE: Uses a mixture of Gaussians (one per digit class), producing sharper samples and a structured latent space with distinct clusters.

## 4 Generative Adversarial Networks (GANs)

### 4.1 Overview

GANs, introduced by Goodfellow et al. (2014), are generative models that train a generator and discriminator in a minimax game to produce realistic data samples.

### 4.2 Core Components

- Generator ( $G$ ): Maps noise  $z \sim p_z(z)$  (e.g.,  $\mathcal{N}(0, I)$ ) to fake samples  $G(z)$ .
- Discriminator ( $D$ ): Outputs probability  $D(x) \in [0, 1]$  to distinguish real data ( $x \sim p_{\text{data}}$ ) from fake ( $G(z)$ ).

### 4.3 Objective Function

$$\min_G \max_D \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log(1 - D(G(z)))]$$

At equilibrium,  $p_g(x) = p_{\text{data}}(x)$ , and  $D(x) = 0.5$ .

### 4.4 Training Process

- Update  $D$  to maximize discrimination accuracy.
- Update  $G$  to minimize  $D(G(z))$ , using non-saturating loss:  $-\mathbb{E}_{z \sim p_z} [\log D(G(z))]$ .

### 4.5 Variants

- DCGANs: Use convolutional layers for image tasks.
- Conditional GANs: Condition on labels or text.
- WGANs: Use Wasserstein distance for stability.
- CycleGAN: Unpaired image-to-image translation.
- StyleGAN: High-resolution image synthesis with style control.

## 4.6 Applications

- Image synthesis (e.g., faces, artwork).
- Image-to-image translation (e.g., sketches to photos).
- Data augmentation, text-to-image synthesis, audio/video generation.

## 4.7 Challenges

- Training instability and mode collapse.
- Difficulty in evaluating sample quality.
- High computational cost and ethical concerns (e.g., deepfakes).

## 4.8 Connection to AAEs

- AAEs use adversarial training to match latent  $q(z)$  to  $p(z)$ , while GANs match generated  $p_g(x)$  to  $p_{\text{data}}(x)$ .
- AAEs combine reconstruction (autoencoder) with generation, unlike GANs, which focus only on generation.

## 4.9 Example: MNIST Digits

- Generator: Maps 100D noise to 28x28 images using dense layers.
- Discriminator: Classifies 28x28 images as real or fake using CNNs.
- Outcome: Generates realistic digits, but lacks reconstruction ability unlike AAEs.